

Assignment 1
Report
Of
Deployment Steps & Implementation
BIG DATA

Team Members

Mostafa Fathi Mohamed
Teacher Assistant @ Nile University

Deployment Steps

1. Docker-compose

- **Namenode Service:**

- ✓ Image: bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8
- ✓ Container Name: namenode
- ✓ Ports: Expose port 9870 on the host, mapped to port 9870 in the container.
- ✓ Volumes: Mount the hadoop_namenode volume to /hadoop/dfs/name in the container.
- ✓ Environment: Set CLUSTER_NAME to test.
- ✓ Environment File: Use the configurations from hadoop.env.

- **Datanode Service:**

- ✓ Image: bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
- ✓ Container Name: datanode
- ✓ Ports: Expose port 9864 on the host, mapped to port 9864 in the container.
- ✓ Volumes: Mount the hadoop_datanode volume to /hadoop/dfs/data in the container.
- ✓ Environment: Set SERVICE_PRECONDITION to "namenode:9870".
- ✓ Environment File: Use the configurations from hadoop.env.

- **Resourcemanager Service:**

- ✓ Image: bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8
- ✓ Container Name: resourcemanager
- ✓ Ports: Expose port 8088 on the host, mapped to port 8088 in the container.
- ✓ Environment: Set SERVICE_PRECONDITION to "namenode:9000 namenode:9870 datanode:9864".
- ✓ Environment File: Use the configurations from hadoop.env.

- **Nodemanager Service:**

- ✓ Image: bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-java8
- ✓ Container Name: nodemanager
- ✓ Ports: Expose port 8042 on the host, mapped to port 8042 in the container.
- ✓ Environment: Set SERVICE_PRECONDITION to "namenode:9000 namenode:9870 datanode:9864 resourcemanager:8088".
- ✓ Environment File: Use the configurations from hadoop.env.

- **Historyserver Service:**

- ✓ Image: bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-java8
- ✓ Container Name: historyserver
- ✓ Ports: Expose port 8188 on the host, mapped to port 8188 in the container.

- ✓ Volumes: Mount the `hadoop_historyserver` volume to `/hadoop/yarn/timeline` in the container.
- ✓ Environment File: Use the configurations from `hadoop.env`.
- **Spark Master Service:**
 - ✓ Image: `bde2020/spark-master:3.0.0-hadoop3.2`
 - ✓ Container Name: `spark-master`
 - ✓ Depends On: `namenode`, `datanode`
 - ✓ Ports: Expose ports 8080 and 7077 on the host, mapped to 8080 and 7077 in the container.
 - ✓ Environment:
 - Set `INIT_DAEMON_STEP` to `setup_spark`.
 - Set `CORE_CONF_fs_defaultFS` to `hdfs://namenode:9000`.
- **Spark Worker 1 Service:**
 - ✓ Image: `bde2020/spark-worker:3.0.0-hadoop3.2`
 - ✓ Container Name: `spark-worker-1`
 - ✓ Depends On: `spark-master`
 - ✓ Ports: Expose port 8081 on the host, mapped to port 8081 in the container.
 - ✓ Environment:
 - Set `SPARK_MASTER` to `spark://spark-master:7077`.
 - Set `CORE_CONF_fs_defaultFS` to `hdfs://namenode:9000`.
- **Volumes:**
 - ✓ Define three named volumes: `hadoop_namenode`, `hadoop_datanode`, and `hadoop_historyserver`. These volumes are used to persist data in the respective containers.

2. Hadoop:

- Download and install Hadoop using `bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8` image in the docker container.
- Configure Hadoop settings in `hadoop-env.sh`, `core-site.xml`, and `hdfs-site.xml`.
- Format the Hadoop Distributed File System (HDFS).
- Start Hadoop services using `start-all.sh` or individual commands (`start-dfs.sh`, `start-yarn.sh`).

3. Spark:

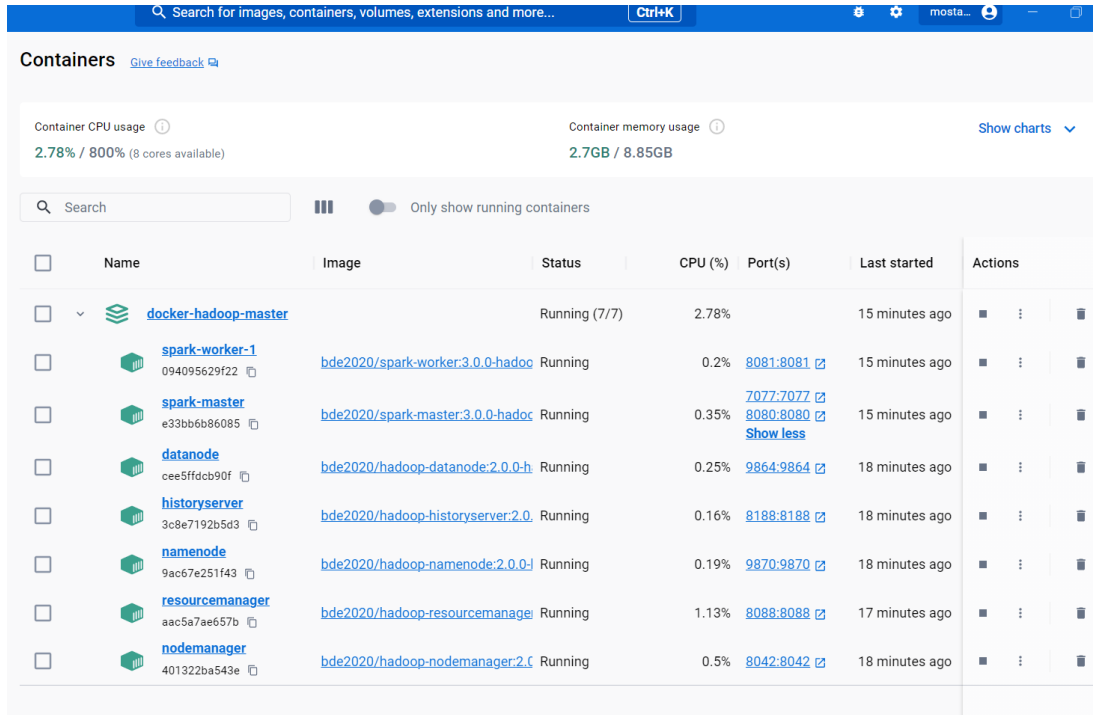
- Download and install Apache Spark using `bde2020/spark-master:3.0.0-hadoop3.2` image in the docker container.
- Configure Spark settings in `spark-env.sh`.
- Set up Hadoop configuration in `spark-defaults.conf`.
- Start Spark services using `start-master.sh` and `start-worker.sh`.

4. YARN:

- Ensure Hadoop is correctly configured and running.
- Configure YARN settings in yarn-site.xml.
- Start the ResourceManager and NodeManagers using start-yarn.sh.

Implementations Steps

Using Docker-Compose to use spark, Hadoop and yarn: -



The screenshot shows the Docker Desktop interface with a list of running containers. The containers are part of a Hadoop and Spark setup. The table below summarizes the data visible in the screenshot:

Name	Image	Status	CPU (%)	Port(s)	Last started
docker-hadoop-master		Running (7/7)	2.78%		15 minutes ago
spark-worker-1	bde2020/spark-worker:3.0.0-hadoop	Running	0.2%	8081:8081	15 minutes ago
spark-master	bde2020/spark-master:3.0.0-hadoop	Running	0.35%	7077:7077, 8080:8080	15 minutes ago
datanode	bde2020/hadoop-datanode:2.0.0-h	Running	0.25%	9864:9864	18 minutes ago
historyserver	bde2020/hadoop-historyserver:2.0	Running	0.16%	8188:8188	18 minutes ago
namenode	bde2020/hadoop-namenode:2.0.0-h	Running	0.19%	9870:9870	18 minutes ago
resourcemanager	bde2020/hadoop-resourcemanager	Running	1.13%	8088:8088	17 minutes ago
nodemanager	bde2020/hadoop-nodemanager:2.0	Running	0.5%	8042:8042	18 minutes ago

■ Hadoop -MapReduce

- Create MapReduce using java to select effective columns and Emit key-value pairs (customer Id as key, selected columns as value)
- Store data in HDFs to be ready for processing from spark.

Now, Data In hdfs after preprocessing: -

```

Command Prompt - docker exec -it namenode bash
C:\Users\Mostafa>docker exec -it namenode bash
root@9ac67e251f43:/# cd /home
root@9ac67e251f43:/home# hdfs dfs -ls /assignment1/output01
Found 2 items
-rw-r--r-- 3 root supergroup 0 2023-12-07 12:11 /assignment1/output01/_SUCCESS
-rw-r--r-- 3 root supergroup 14623 2023-12-07 12:11 /assignment1/output01/part-r-00000
root@9ac67e251f43:/home# hdfs dfs -cat /assignment1/output01/part-r-00000
2023-12-09 16:34:00,429 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
101 Female,29,New York,1120.2,14,Satisfied
102 Male,34,Los Angeles,780.5,11,Neutral
103 Female,43,Chicago,510.75,9,Unsatisfied
104 Male,30,San Francisco,1480.3,19,Satisfied
105 Male,27,Miami,720.4,13,Unsatisfied
106 Female,37,Houston,440.8,8,Neutral
107 Female,31,New York,1150.6,15,Satisfied
108 Male,35,Los Angeles,800.9,12,Neutral
109 Female,41,Chicago,495.25,10,Unsatisfied
110 Male,28,San Francisco,1520.1,21,Satisfied
111 Male,32,Miami,690.3,11,Unsatisfied
112 Female,36,Houston,470.5,7,Neutral
113 Female,30,New York,1200.8,16,Satisfied
114 Male,33,Los Angeles,820.75,13,Satisfied
115 Female,42,Chicago,530.4,9,Unsatisfied
116 Male,29,San Francisco,1360.2,18,Satisfied
117 Male,26,Miami,700.6,12,Unsatisfied
118 Female,38,Houston,450.9,8,Neutral
119 Female,32,New York,1170.3,14,Satisfied
120 Male,34,Los Angeles,700.2,11,Neutral

```

■ Spark

- Start execute spark-master container and launch spark session: -
 - 1- `docker exec -it spark-master bash`
 - 2- `/spark/bin/pyspark --master spark://spark-master:7077`
- Determine data path from hdfs: -


```
data_path = "hdfs://namenode:9000/assignment1/output01/part-r-00000"
```
- Convert Our data To RDD to start to work in it:-


```
data_rdd = spark.sparkContext.textFile(data_path)
```

Apply some of data processing to reformat data from (key, value) to be data frame like this: -



Spark Master at spark://e33bb6b86085-7077

URL: spark://e33bb6b86085:7077
Alive Workers: 1
Cores in use: 8 Total, 8 Used
Memory in use: 8.1 GiB Total, 1024.0 MiB Used
Resources in use:
Applications: 1 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (1)

Worker Id
worker-20231209155404-172.18.0.8-37047

Running Applications (1)

Application ID	Name	Cores
app-20231209160057-0000	(kill) PySparkShell	8

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor
----------------	------	-------	---------------------

Command Prompt - docker exec -it spark-master bash


```
NameError: name 'data_path' is not defined
>>> data_path = "hdfs://namenode:9000/assignment1/output01/part-r-00000"
>>> data_rdd = spark.sparkContext.textFile(data_path)
>>> header = data_rdd.take(1)[0]
>>> data_rows = data_rdd.filter(lambda line: line != header)
>>> formatted_data = data_rows.map(lambda x: x.replace('\t', ','))
>>> split_data = formatted_data.map(lambda line: line.split(","))
>>> schema = StructType([
...     StructField("Customer ID", StringType(), True),
...     StructField("Gender", StringType(), True),
...     StructField("Age", StringType(), True),
...     StructField("City", StringType(), True),
...     StructField("Total Spend", StringType(), True),
...     StructField("Items Purchased", StringType(), True),
...     StructField("Satisfaction Level", StringType(), True)
... ])
>>> df = spark.createDataFrame(split_data, schema=schema)
>>> df.show()
```

Customer ID	Gender	Age	City	Total Spend	Items Purchased	Satisfaction Level
102	Male	34	Los Angeles	780.5	11	Neutral
103	Female	43	Chicago	510.75	9	Unsatisfied
104	Male	30	San Francisco	1480.3	19	Satisfied
105	Male	27	Miami	720.4	13	Unsatisfied
106	Female	37	Houston	440.8	8	Neutral
107	Female	31	New York	1150.6	15	Satisfied
108	Male	35	Los Angeles	800.9	12	Neutral
109	Female	41	Chicago	495.25	10	Unsatisfied
110	Male	28	San Francisco	1520.1	21	Satisfied

- After Applying Preprocessing and some of Data Engineering using spark, we stored output in hdfs:-
output_path = "hdfs://namenode:9000/assignment1/output_csv"
csv_data = split_data.map(lambda line: ",".join(line))
csv_data.saveAsTextFile(output_path)
- We are using collab to apply spark ml and use the output from the previous operation which is data stored in HDFS.
- Now, our data contain the following columns: -

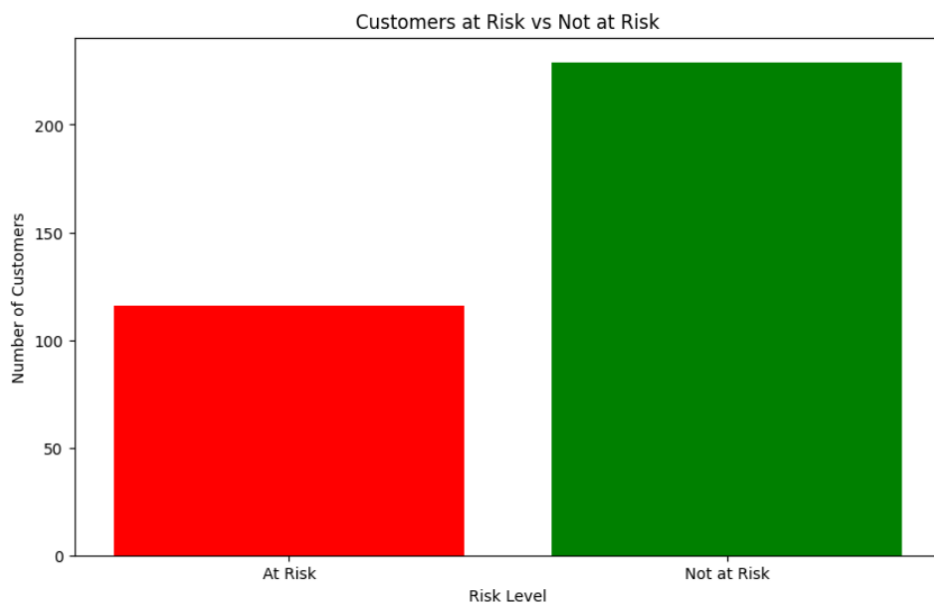
```
>>> df.printSchema()
root
 |-- Customer ID: string (nullable = true)
 |-- Gender: string (nullable = true)
 |-- Age: string (nullable = true)
 |-- City: string (nullable = true)
 |-- Total Spend: string (nullable = true)
 |-- Items Purchased: string (nullable = true)
 |-- Satisfaction Level: string (nullable = true)
```

- Start our work to answer 3 Questions: -
 - Can we segment customers based on their demographic information (Age, Gender, City) and shopping behaviors (Total Spend, Number of Items Purchased, Membership Type)?
➔ To answer this question, we applied KMeans algorithm and using scatter plot to visualize our results.



2. Which customers are at risk of not making future purchases based on their Days?

➔ To answer this question, we used RandomForestClassifier to classify customers with some features specially "Days Since Last Purchase", "Items Purchased" and "Total Spend"



3. Can we predict a customer's Satisfaction Level based on their demographic and purchase history data?

➔ To answer this question, we use DecisionTreeClassifier to classify our customers by using feature like "Age", "Total Spend", "Items Purchased", "GenderIndex", "CityIndex" and use MulticlassClassificationEvaluator to evaluate our model and test model by trying to predict new data and get result as following:-


```
new_data_samples = [  
    (300, "Male", 33, "Chicago", 800.5, 14),  
    (301, "Female", 29, "New York", 1200.7, 17),  
    (302, "Male", 35, "Los Angeles", 700.8, 11)  
]
```

✓ The output

```
+-----+-----+-----+  
|Customer ID|          features|prediction|  
+-----+-----+-----+  
|      300|[33.0,800.5,14.0,...|      2.0|  
|      301|[29.0,1200.7,17.0...|      0.0|  
|      302|[35.0,700.8,11.0,...|      1.0|  
+-----+-----+-----+
```

Yarn

- Using as resource management layer for Apache Hadoop allows Spark and other applications to efficiently share and allocate resources in a Hadoop cluster. And the following is screenshot from resource manager for all application and screenshot for specific application: -



All Applications

Cluster

About Nodes Node Labels Applications

NEW NEW SAVING SUBMITTED ACCEPTED RUNNING FINISHED FAILED KILLED

Scheduler

Tools

Cluster Metrics															
Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved					
0	0	0	0	0	0 B	16 GB	0 B	0	8	0					

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=M), vcores]	<memory 1024, vCores 1>	<memory 8192, vCores 4>	0

Show 20 entries

ID	User	Name	Application Type	Queues	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Reserved CPU Vcores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1701950512249_0001	root	ECommerceMapReduce	MAPREDUCE	default	0	Thu Dec 7 14:11:05 +0200 2023	Thu Dec 7 14:11:07 +0200 2023	Thu Dec 7 14:11:37 +0200 2023	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0009	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 23:51:40 +0200 2023	Mon Dec 4 23:51:41 +0200 2023	Mon Dec 4 23:52:03 +0200 2023	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0008	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 23:29:32 +0200 2023	Mon Dec 4 23:29:33 +0200 2023	Mon Dec 4 23:30:06 +0200 2023	FINISHED	FAILED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0007	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 23:01:05 +0200 2023	Mon Dec 4 23:01:08 +0200 2023	Mon Dec 4 23:01:29 +0200 2023	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0006	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 22:56:42 +0200 2023	Mon Dec 4 22:56:43 +0200 2023	Mon Dec 4 22:57:04 +0200 2023	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0005	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 22:47:35 +0200 2023	Mon Dec 4 22:47:36 +0200 2023	Mon Dec 4 22:48:01 +0200 2023	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0004	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 22:43:21 +0200 2023	Mon Dec 4 22:43:21 +0200 2023	Mon Dec 4 22:43:55 +0200 2023	FINISHED	FAILED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0
application_1701721101940_0003	root	ECommerceMapReduce	MAPREDUCE	default	0	Mon Dec 4 22:43:21 +0200 2023	Mon Dec 4 22:43:21 +0200 2023	Mon Dec 4 22:43:55 +0200 2023	FINISHED	FAILED	N/A	N/A	N/A	N/A	N/A	0.0	0.0		History	0



Application application_1701950612249_0001

Logged in as: root

Cluster

About
Nodes
Node Labels
Applications
NEW
NEW SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED
Scheduler

Tools

Application Overview

User: root
Name: ECommerceMapReduce
Application Type: MAPREDUCE
Application Tags:
Application Priority: 0 (Higher Integer value indicates higher priority)
YarnApplicationState: FINISHED
Queue: default
FinalStatus Reported by AM: SUCCEEDED
Started: Thu Dec 07 12:11:05 +0000 2023
Launched: Thu Dec 07 12:11:07 +0000 2023
Finished: Thu Dec 07 12:11:37 +0000 2023
Elapsed: 31sec
Tracking URL: History
Log Aggregation Status: NOT_START
Application Timeout (Remaining Time): Unlimited
Diagnostics: Attempt recovered after RM restart
Unmanaged Application: false
Application Node Label expression: <Not set>
AM container Node Label expression: <DEFAULT_PARTITION>

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>
Total Number of Non-AM Containers Preempted: 0
Total Number of AM Containers Preempted: 0
Resource Preempted from Current Attempt: <memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt: 0
Aggregate Resource Allocation: 147508 MB-seconds, 43 vcore-seconds
Aggregate Preempted Resource Allocation: 147508 MB-seconds, 43 vcore-seconds

Show 20 entries

Search

Attempt ID	Started	Node	Logs	Nodes blacklisted by the app	Nodes blacklisted by the system
appattempt_1701950612249_0001_000001	Thu Dec 7 14:11:06 +0200 2023	http://54f57a5f37d5.8042	Logs	0	0

Showing 1 to 1 of 1 entries

First Previous 1 Next Last